

Pretrained model
Qwen3-8B (inference)

① LLM-based translation generation (GPT-4o)

Input: ACL60/60 dev source sentences, allowed actions

Output: action sequence + translations



② Per-action statistics aggregation

Input: action-based translations

Output: action-level quality and latency statistics



③ Prompt construction



④ Step-wise inference (Qwen3-8B)

Input: source sentence + prompt

Output: incremental translation



⑤ Evaluation

ACL60/60 eval (all methods); MUST-C tst-HE (mBART)

Metrics: BLEU, chrF, TER, COMET, LAAL